Workshop:
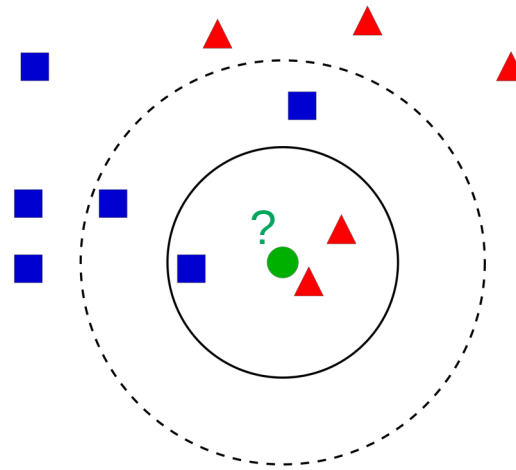**Machine Learning and Prediction Modelling**

K Nearest Neighbor

Yannick Rothacher

SPF, HS2025

Recap: k-means clustering

▶ The k-means algorithm is a **partitional clustering** method

▶ Often there is not a "correct" solution in k-means clustering

   ▶ The goal is to find out whether there is a cluster-structure present in the data (find hidden patterns)

▶ Such an approach is generally referred to as **unsupervised** learning

▶ In contrast, in **supervised** learning we try to predict the correct value of an observed variable

   ▶ Techniques, which try to predict the correct class of data are referred to as **classifiers**
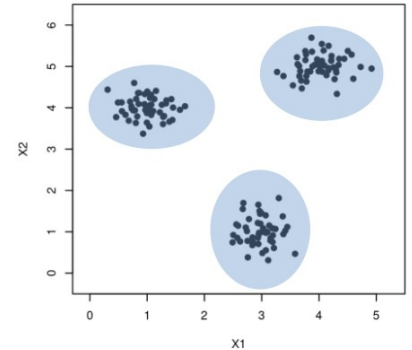
Supervised learning: Classification

▸ Classifiers can be used to predict all kinds of data labels

    ▸ Who will win in the next football-worldcup?

    ▸ Will an operation be successful?

    ▸ Which type of tumor is shown on a medical scan?

▸ Predictions are always based on training data:

Hours studied	Score last year	IQ	Exam result
1	20	101	**failed**
20	99	105	**passed**
12	54	99	**passed**
...	...	...	...

Target variable

New observation:
Will this person pass or fail?

| Hours studied	Score last year	IQ	Exam result
25	70	98	**?**

Supervised learning: Classification

▶ The workflow for classifiers is always similar

▶ First, a model is fitted to **training** data, for which the true class labels are known

▶ Afterwards, the model is applied to new observations to predict their class labels

▶ Thus, we always first need training data to **train our model**

▶ There are different statistical models, which can be used as a classifier

  ▶ e.g. **logistic regression** is a simple classifier usually applied when the target label has two levels (e.g. predict sex: male/female)

  ▶ **K-nearest neighbor** is an easy to understand classifier and, therefore, a good introduction to the topic of classification

Hours studied	Score last year	IQ	Exam result
1	20	101	failed
20	99	105	passed
12	54	99	passed
...	...	...	...

K-nearest neighbor (KNN)

► Not to confuse with k-means algorithm!

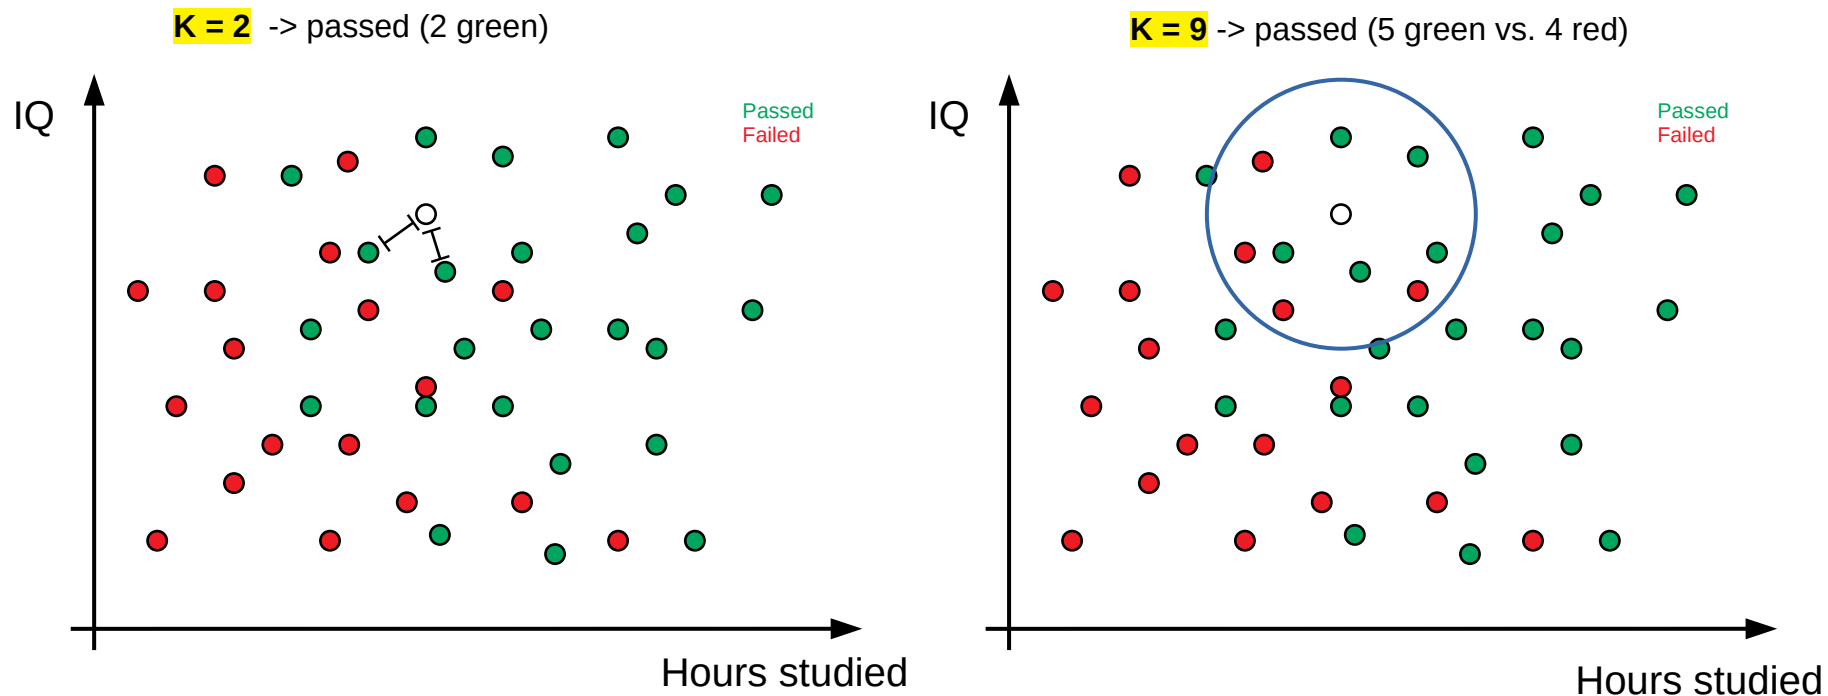► Idea: Classify a new data-point based on the labels of the neighboring training data points

▷ Example in 2D:
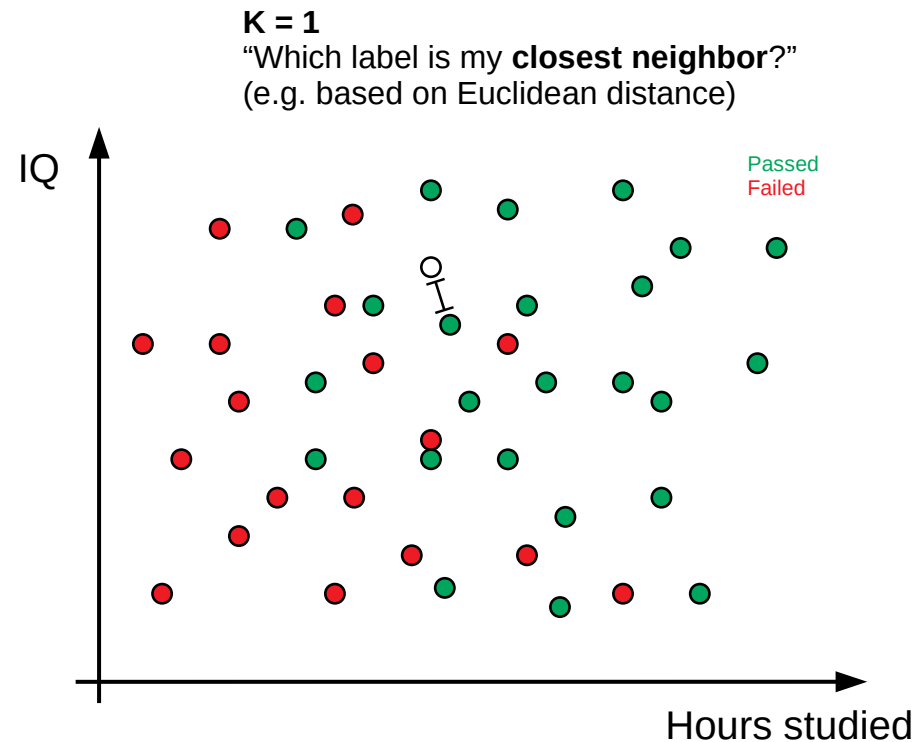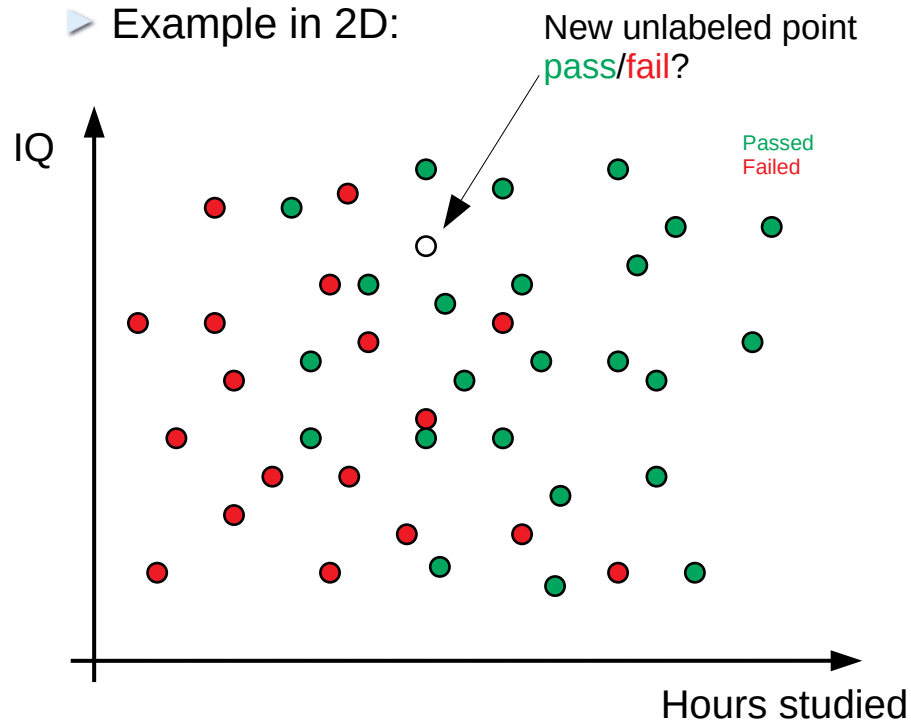
► Assign the label "passed" because closest neighbor is "passed" as well

K-nearest neighbor (KNN)

▶ The parameter k determines how many neighbors are considered (use the majority vote as prediction):

K = 2 -> passed (2 green)

K = 9 -> passed (5 green vs. 4 red)

▶ The prediction may change depending on k

▶ There is the possibility of ties, **what could be done to deal with ties?**
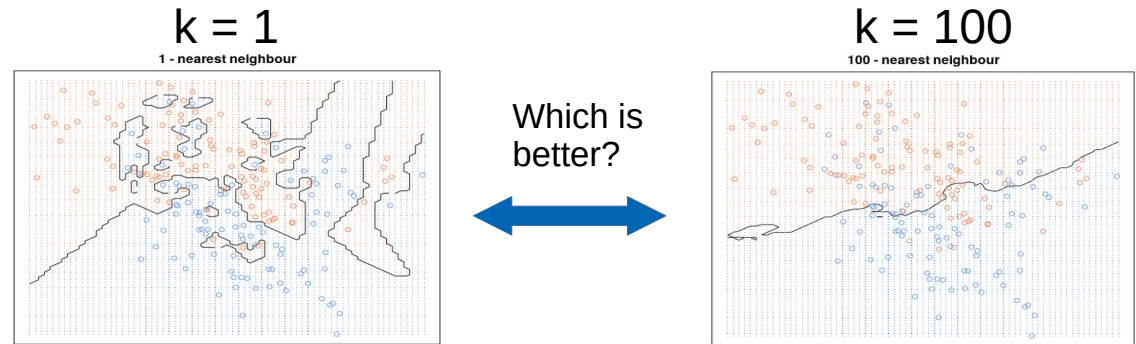
KNN - Decision boundaries

▶ The decision boundaries for new data points depend on how many neighbors are considered:

▶ The decision boundary becomes **simpler** with larger k

▶ What is the **best** value of k?

How to choose k

► What is the best value of k?

k = 1
1 - nearest neighbour

Which is better?

k = 100
100 - nearest neighbour

► In order to decide what the best value for k is we want to **measure** how "well" the classifier is performing

  ► How can we do that?

► Simple idea: Check how correctly the classifier is labeling the training data (**Training error**)

  ► Confusion matrices (for k=1 and k=100):

k = 1

```
                  true label
KNN prediction  blue red
          blue   100    0
          red      0  100
```

k = 100

```
                  true label
KNN prediction  blue red
          blue    66   26
          red     34   74
```

Train- and test-error

▶ Calculate the training error (number of miss-classifications divided by total number of predictions):

<div align="center">

k = 1

true label

KNN prediction blue red

blue 100   0

red   0 100

→ 0/200 = **0%**

</div>

<div align="center">

k = 100

true label

KNN prediction blue red

blue   66   26

red   34   74

→ 60/200 = **30%**

</div>

▶ Looking at the training error is not that informative. In our example, K=1 forces a 0% training error rate.

▶ Because we want to use the classifier to predict the class of new observations (not included in the training set), it makes sense to calculate the error rate for unseen data (**Test-error**)
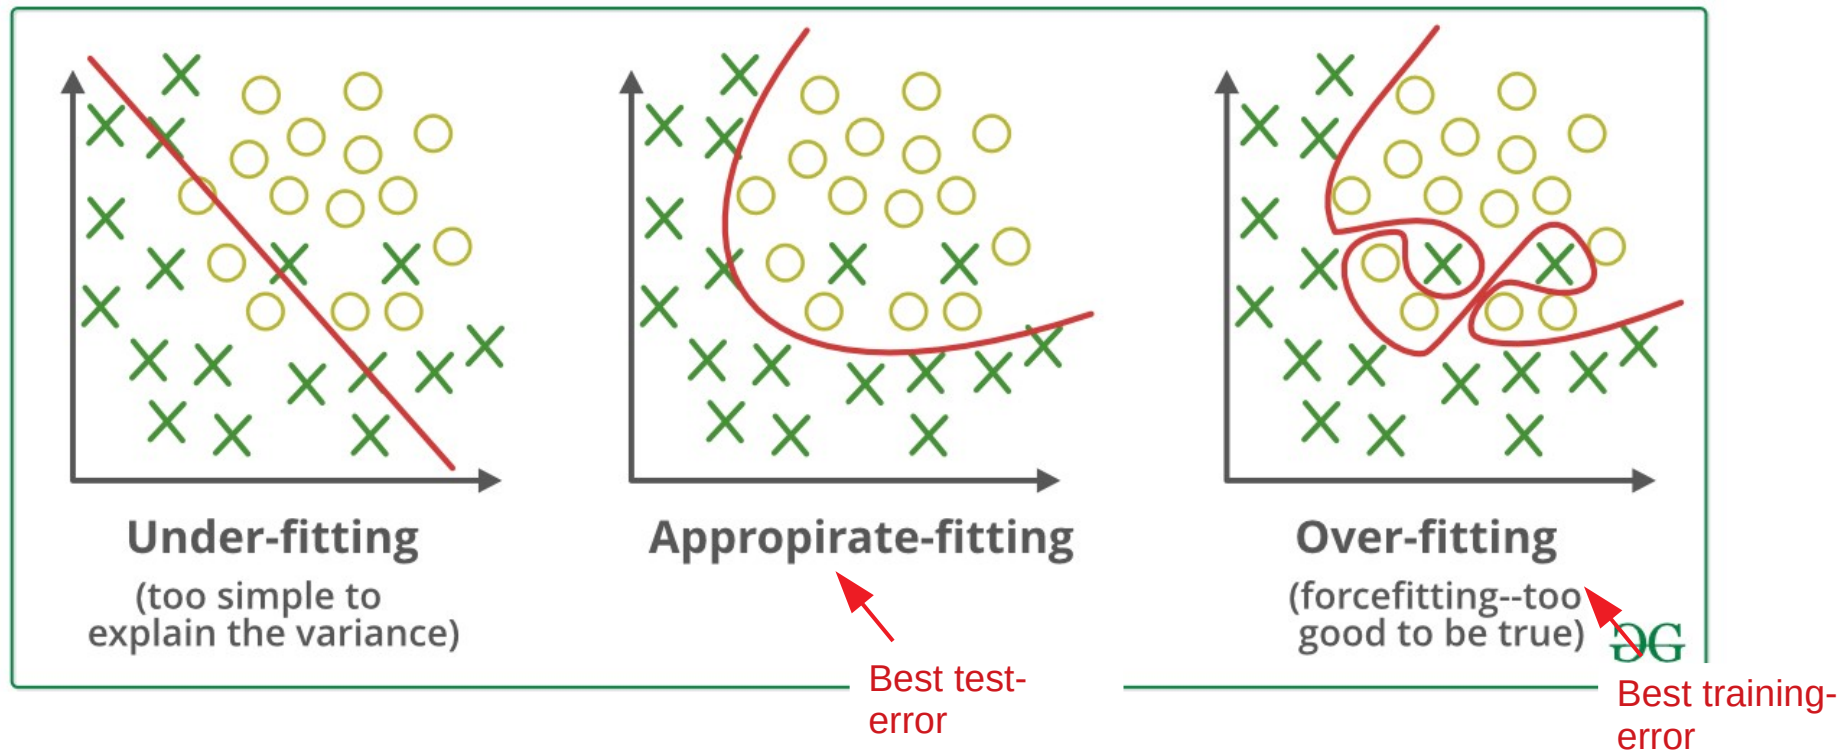
Train- and test-error

▶ **Test-error** for 40 new observations (for k=1 and k=100)

<div>

k = 1

true label

KNN prediction blue red

blue   12   8

red   3 17

→ 11/40 = **27.5%**

k = 100

true label

KNN prediction blue red

blue   13   7

red   2 18

→ 9/40 = **22.5%**

</div>

▶ Now the two classifiers both perform equally well (in this case k=100 even slightly better)

▶ The goal is to keep the **test-error** as low as possible

▶ Most probably, setting k=1 and setting k=100 are both not ideal

  ▷ k=1 is too flexible (**over-fitting** the training data) and k=100 is perhaps too simple (**under-fitting** the training data)
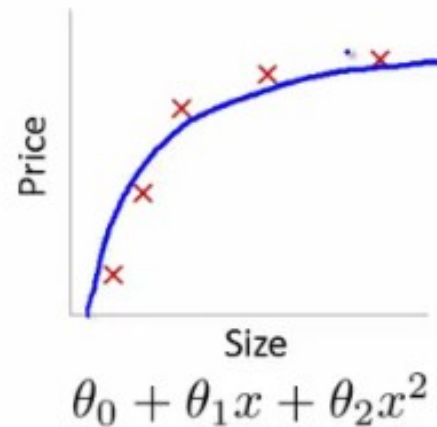
Over- vs. under-fitting

Under-fitting
(too simple to explain the variance)

Appropirate-fitting

Over-fitting
(forcefitting--too good to be true)

Best test-error

Best training-error

▶ Imagine you apply repeatedly an over- or under-fitting model to collected data:

  ▶ Under-fitting will show **high bias** and **low variance**

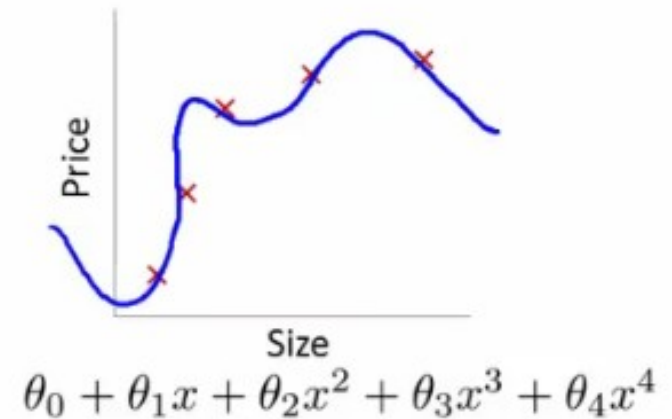 ▶ Over-fitting will show **low bias** and **high variance**

Over- vs. under-fitting (in regression)

$$\theta_0 + \theta_1 x$$

High bias (underfit)

$$\theta_0 + \theta_1 x + \theta_2 x^2$$

"Just right"

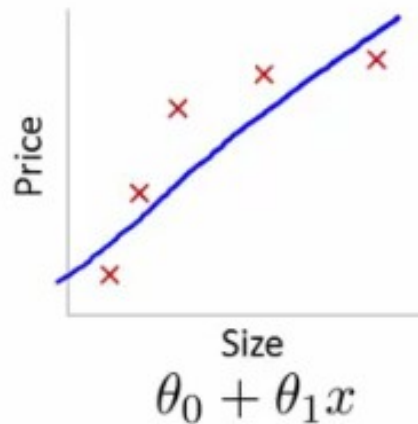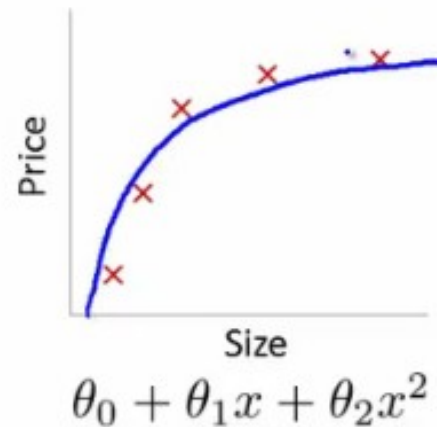$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$
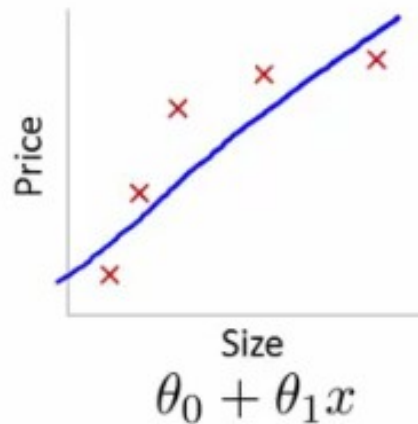
High variance (overfit)

▶ Imagine you apply repeatedly an over- or under-fitting model to collected data:

  ▷ Under-fitting will show **high bias** and **low variance**

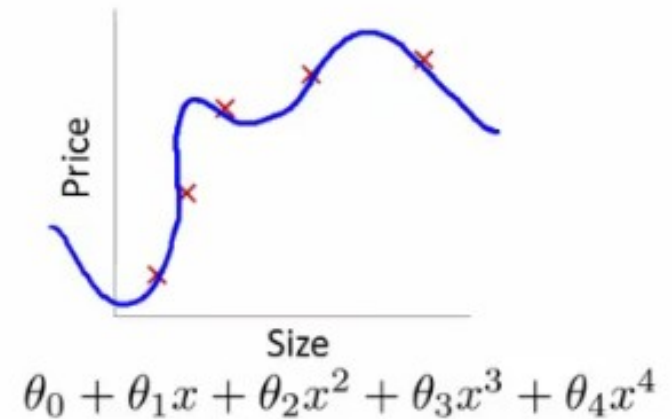 ▷ Over-fitting will show **low bias** and **high variance**

Over- vs. under-fitting (in regression)

High bias (underfit): $\theta_0 + \theta_1 x$

"Just right": $\theta_0 + \theta_1 x + \theta_2 x^2$

High variance (overfit): $\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$

▶ Imagine you apply repeatedly an over- or under-fitting model to collected data:

 ▷ Under-fitting will show **high bias** and **low variance**

 ▷ Over-fitting will show **low bias** and **high variance**

→ show in R

https://towardsdatascience.com/underfitting-and-overfitting-in-machine-learning-and-how-to-deal-with-it-6fe4a8a49dbf

Where is the optimum between over- and under-fitting?

► Test-error should be as small as possible
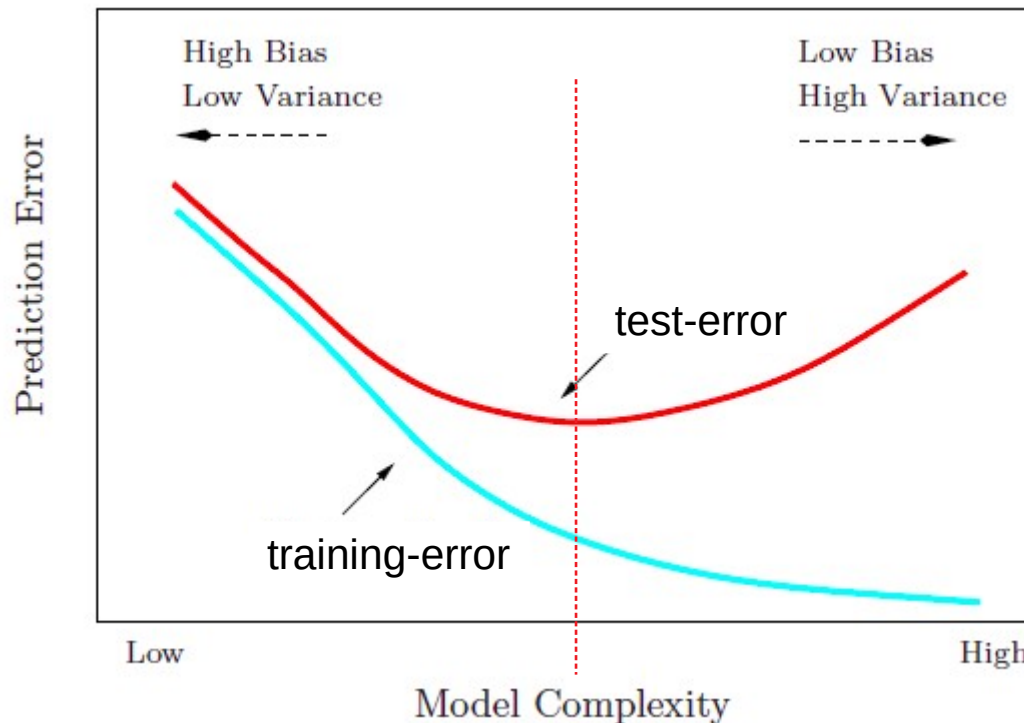
► Dashed, red line shows the **ideal model flexibility**

  ► Everything to the right of dashed line is over-fitting

  ► Everything to the left of dashed line is under-fitting

Outlook: Cross-validation

- It is important to **estimate the test-error** of a classifier to judge its performance

- Ideally, one uses a **training data set** to fit the model and separate **test data set** to assess its performance

    - Often a separate test data set is not available

- In the next topic we see what methods are used to estimate the test-error as efficient as possible (cross-validation)

KNN in R

```r
library(class)
knn.pred <- knn(train = iris[,1:4],   # use the iris data as training data
                cl = iris$Species,  # correct labels of training data
                test = iris[,1:4],   # knn always needs a specified test set
                k = 5)                  # parameter k
# Confusion matrix:
confT <- table(knn.pred, iris$Species)
confT

knn.pred     setosa versicolor virginica
  setosa         50          0         0
  versicolor      0         47         2
  Virginica       0          3        48

# Training error:
missT <- confT
diag(missT) <- 0  # turn diagonal values to zeros
missCount <- sum(missT)    # count the miss-classifications (5)
trainErr <- missCount/sum(confT)
trainErr
[1] 0.03333333
```

Because we do not have a test set, we reuse the training data as test data

Standardization in KNN

- Similar to the case of PCA it often makes sense to standardize the data before applying KNN classification

  - Why could standardization be important?

- **Example**: If one variable has a much larger variance than all the others (e.g. ranging from 0 - 100'000, other variables: 0 - 0.1), it will be much more important in determining the Euclidean distance between two data points.

  - Other variables are neglected!

- Through standardization, all variables have equal variance

 - In R: use the scale() function to standardize numeric vectors
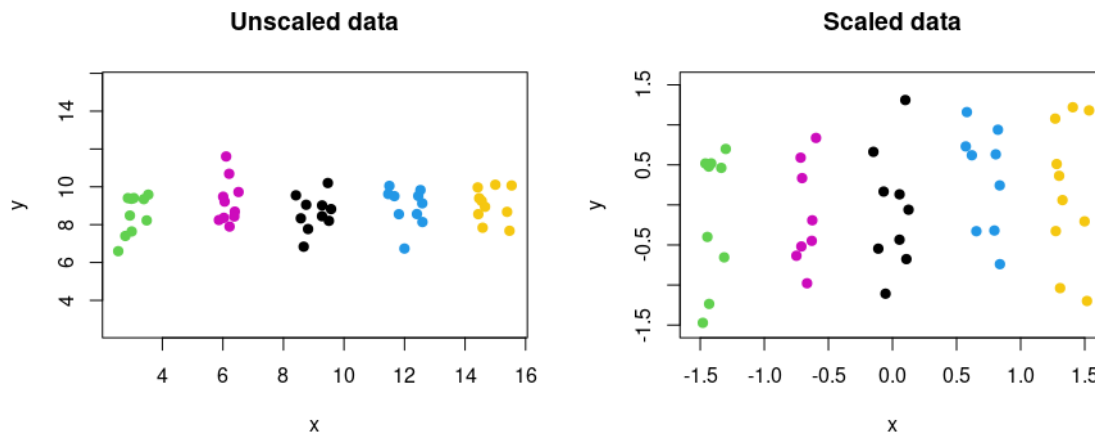
KNN for regression

▷ KNN cannot only be used for classification

    ▷ Can also predict **numerical values**

    ▷ Instead of taking the majority vote of neighbors, take e.g. the **mean of k neighbors**
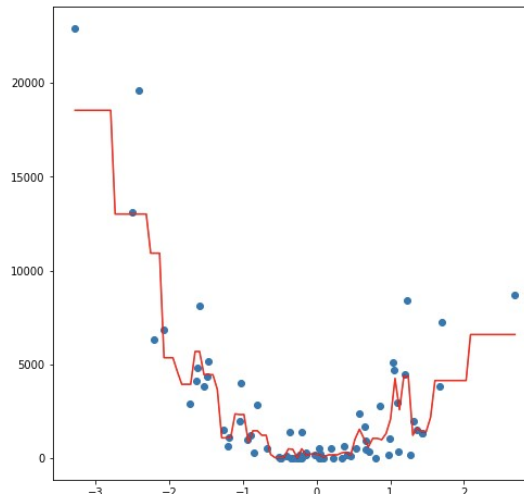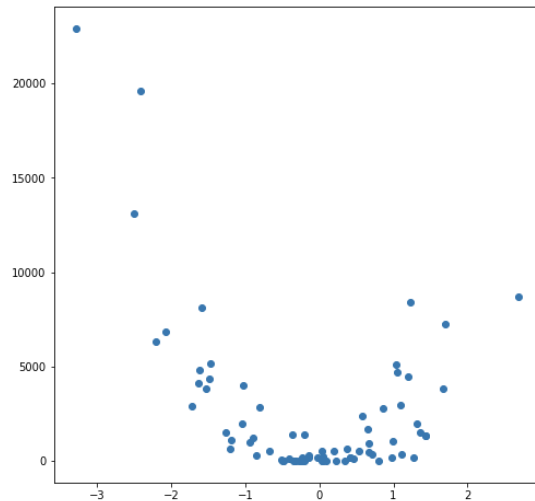

New point:
x = -0.8 y = ?

Predicted value
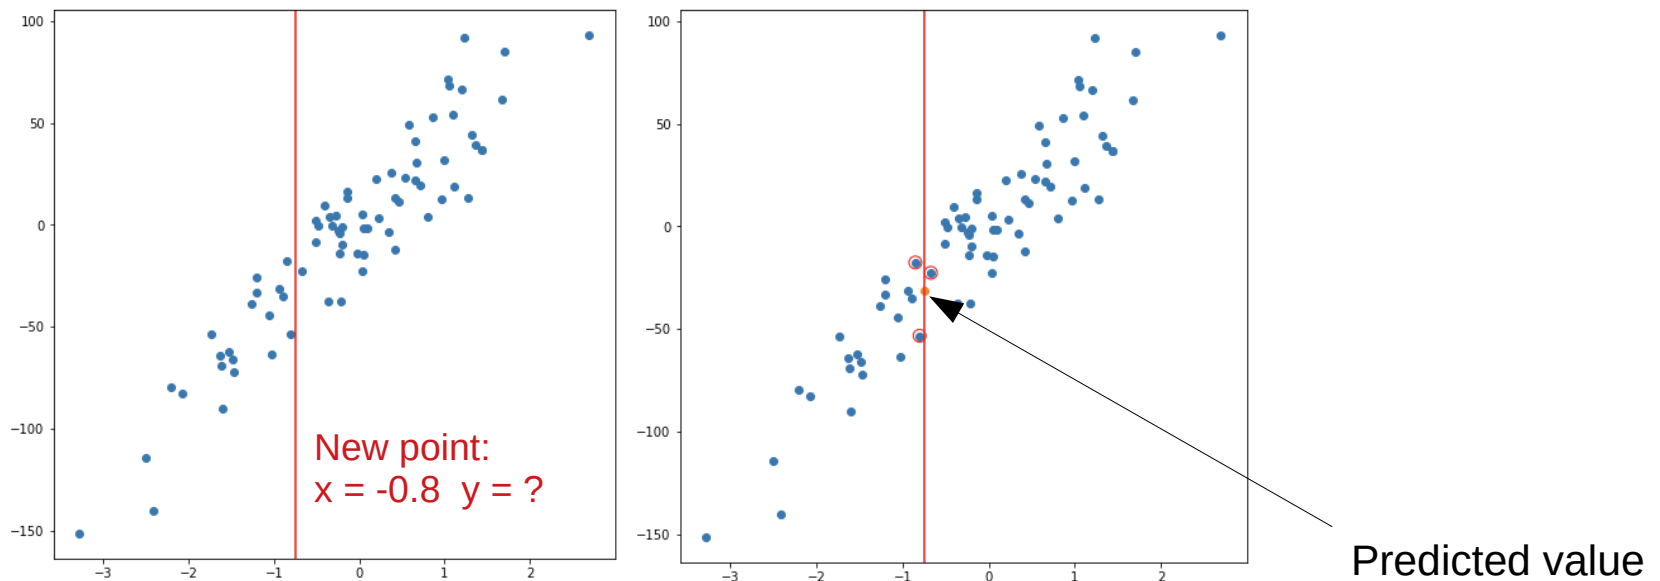
▷ E.g. univariate data (figures above): Use the average of **k=3** closest neighbors (with respect to x value) as a prediction for y value

    ▷ Can use the distances as **weights** for the mean

KNN for regression
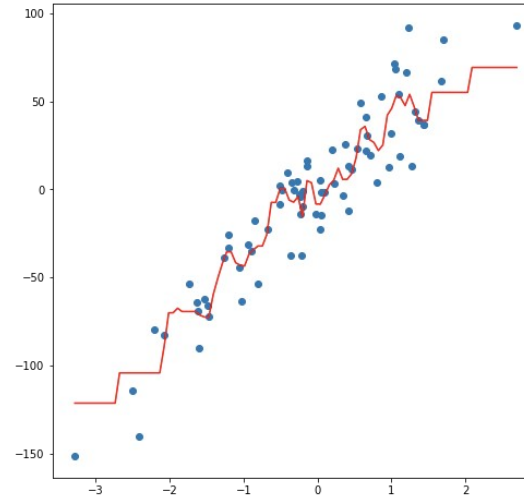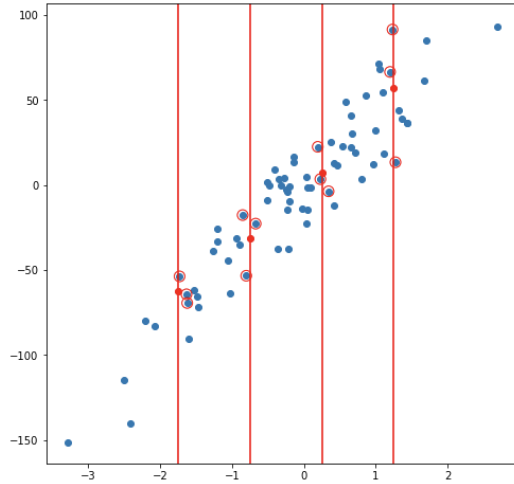
- ▶ KNN for linear relation

- ▶ Result is not a clean line but still general pattern is captured

- ▶ Advantage of KNN for regression:

  Simple algorithm which can capture non-linearity quite well