


*Workshop:*  
**Machine Learning and Prediction Modelling**

# Neural Networks



Yannick Rothacher

*SPF, HS2025*

# Artificial Neural Networks (today)

- ▶ (Deep) Neural Networks are very popular!
  - ▶ Also referred to as "Deep Learning"
- ▶ Successfully applied in many fields:

## Natural language processing




Speech recognition



Large language models (chatbots)

## Reinforcement learning

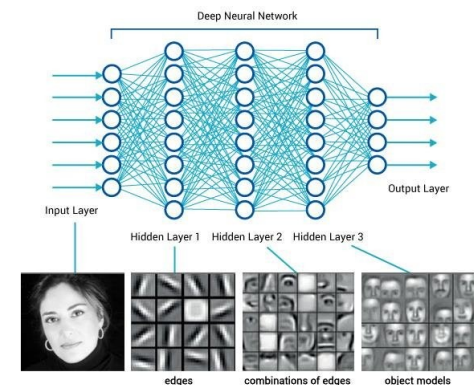


Autonomous driving



Playing games (e.g. alpha go)

## Computer vision



Picture recognition




Image generation

# Artificial Neural Networks (yesterday)

- ▶ **Heureka** exhibition in Zürich, Brunau (**1991**)
- ▶ Presented the "Forschungsstandort Schweiz"



[https://www.e-pics.ethz.ch/index/ethbib.bildarchiv/ETHBIB.Bildarchiv\\_Com\\_FC24-8002-0196\\_24364.html](https://www.e-pics.ethz.ch/index/ethbib.bildarchiv/ETHBIB.Bildarchiv_Com_FC24-8002-0196_24364.html)



## 2 Einzelworterkennung mit neuronalen Netzen (6.3.1)




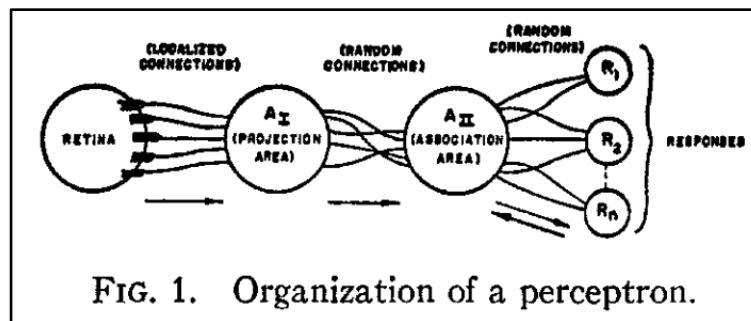
Neuronale Netze werden seit kurzem in einer Vielzahl verschiedener Gebiete verwendet: Bildverarbeitung, Signalbereinigung, Trendanalyse, Regelungstechnik usw. Daneben werden sie weltweit auf

Daneben werden sie weltweit auf ihre Tauglichkeit zur Erkennung gesprochener Sprache untersucht. Das Problem dabei ist, dass die Erkennung ganzer Wörter sprecherunabhängig sein soll. Als Beispiel suchen wir in einem Computer ein Dokument anstatt mit einer Maus mittels gesprochener Schlüsselwörter.




# Artificial Neural Networks (yesterday)

- ▶ The idea of artificial neural networks is very old
  - ▶ First reference dates back to **1944** (Warren S. McCulloch and Walter Pitts)
- ▶ The "**perceptron**" (the first "modern" neural network)
  - ▶ Invented by psychologist **F. Rosenblatt (1958)**

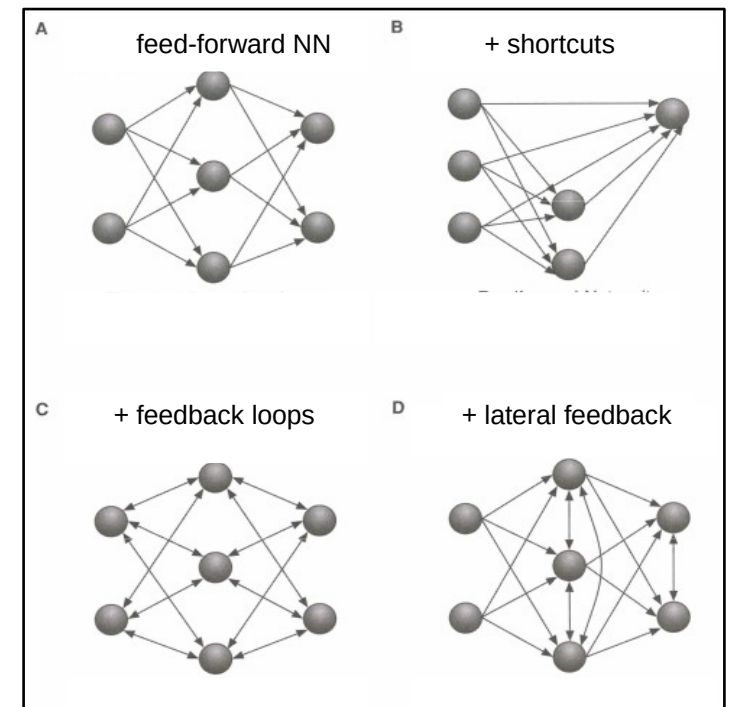


# Structure of Neural Networks (NN)

- Usually represented as connected **nodes** (neurons) organized in **layers**
  - Different structures are possible
- We will focus on **feed-forward** networks




- Structure consists of **input**, **hidden** and **output** layer(s)
- Connections are associated with specific **weights** (strength of connection)
  - Nodes are associated with specific **activation functions** (later)
- The structure is loosely inspired by real-life neurons



Source: Holling & Schmitz (2010)

# Structure of Neural Networks (NN)


- Usually represented as connected **nodes** (neurons) organized in **layers**
  - Different structures are possible
- We will focus on **feed-forward** networks



Input - Layer

Hidden - Layer(s)

Output - Layer



Source: Huang et al. (2018)

- Structure consists of **input**, **hidden** and **output** layer(s)
- Connections are associated with specific **weights** (strength of connection)
  - Nodes are associated with specific **activation functions** (later)
- The structure is loosely inspired by real-life neurons

# Single hidden layer NN


- ▶ Supervised learning with NNs (predict target variable)

- ▶ Exemplary regression data:

Target variable (y)

Training time (X1)	Sleep time (X2)	Body height (X3)	Performance (0-100)
2	8	1.8	60
8	9	1.6	100
5	4	1.7	85
8	5	1.8	79
7	7	1.8	62
...	...	...	...

- ▶ Applied (single hidden layer) NN:



**Input**

Size = 3  
(3 predictors)

**Hidden**

Size = 2  
(free to choose)

**Output**

Size = 1  
(y is numeric)

$\alpha_1 - \alpha_8$  : connection weights

# Single hidden layer NN

- ▶ Supervised learning with NNs (predict target variable)

- ▶ Exemplary regression data:


Training time (X1)	Sleep time (X2)	Body height (X3)	Performance (0-100)
2	8	1.8	60
8	9	1.6	100
5	4	1.7	85
8	5	1.8	79
7	7	1.8	62
...	...	...	...

Target variable (y)

- ▶ The goal is to find the weights, which allow the prediction of **y**

- ▶ How the prediction works (**slightly simplified**, see later):

- ▶ Applied (single hidden layer) NN:



**Input**

Size = 3  
(3 predictors)

**Hidden**

Size = 2  
(free to choose)

**Output**

Size = 1  
(y is numeric)

$\alpha_1 - \alpha_8$ : connection weights

X1	X2	X3	Performance (0-100)
3	5	1.7	?

Value at **H1**:  $H1 = \alpha_1 * 3 + \alpha_3 * 5 + \alpha_5 * 1.7$

Value at **H2**:  $H2 = \alpha_2 * 3 + \alpha_4 * 5 + \alpha_6 * 1.7$

Value at **O1** (= Prediction):  $O1 = \alpha_7 * H1 + \alpha_8 * H2$



# Single hidden layer NN

- ▶ Supervised learning with NNs (predict target variable)

- ▶ Exemplary regression data:

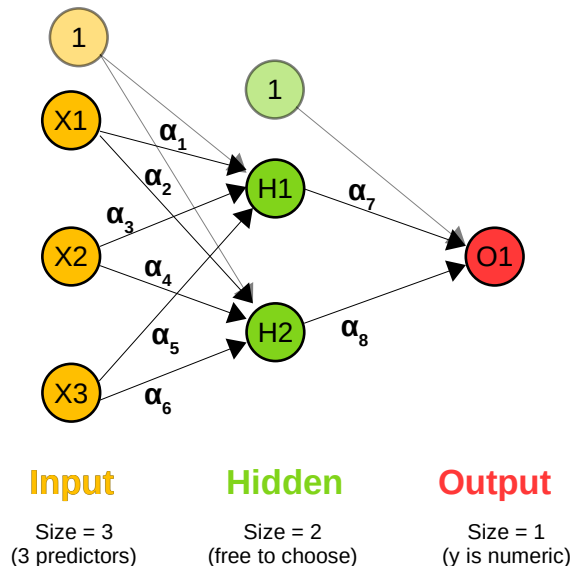
Training time (X1)	Sleep time (X2)	Body height (X3)	Performance (0-100)
2	8	1.8	60
8	9	1.6	100
5	4	1.7	85
8	5	1.8	79
7	7	1.8	62
...	...	...	...

Target variable (y)

- ▶ The goal is to find the weights, which allow the prediction of  $y$

- ▶ How the prediction works (**slightly simplified**, see later):

- ▶ Applied (single hidden layer) NN:



$\alpha_1 - \alpha_8$ : connection weights

X1	X2	X3	Performance (0-100)
3	5	1.7	?

Value at **H1**:  $H1 = \alpha_1 * 3 + \alpha_3 * 5 + \alpha_5 * 1.7$

Value at **H2**:  $H2 = \alpha_2 * 3 + \alpha_4 * 5 + \alpha_6 * 1.7$

Value at **O1** (= Prediction):  $O1 = \alpha_7 * H1 + \alpha_8 * H2$



- ▶ "Biases" are often added at each layer (think of them as "intercepts")

# Prediction of NN in more detail

- ▶ Univariate regression: Only one predictor ( $x$ )

- ▶  $y = b_0 + b_1x + \varepsilon \quad (\varepsilon \sim N(0, \sigma^2))$

- ▶ Applied Neural Network:





Truth:  
 $b_0 = 2$   
 $b_1 = 4$

# Prediction of NN in more detail

- ▶ Univariate regression: Only one predictor ( $x$ )

- ▶  $y = b_0 + b_1x + \varepsilon \quad (\varepsilon \sim N(0, \sigma^2))$

- ▶ Applied Neural Network:




Truth:  
 $b_0 = 2$   
 $b_1 = 4$

# Prediction of NN in more detail

- Univariate regression: Only one predictor (x)

- $y = b_0 + b_1x + \varepsilon \quad (\varepsilon \sim N(0, \sigma^2))$

- Applied Neural Network:



- Prediction of y:

$$H = \alpha_1 + \alpha_2 * x$$


$$y = \alpha_3 * H$$

→ Ideal solution:

$$\alpha_1 = b_0 = 2$$

$$\alpha_2 = b_1 = 4$$

$$\alpha_3 = 1$$




Truth:  
 $b_0 = 2$   
 $b_1 = 4$

# Prediction of NN in more detail

- Univariate regression: Only one predictor (x)

- $y = b_0 + b_1x + \varepsilon \quad (\varepsilon \sim N(0, \sigma^2))$

- Applied Neural Network:



- Prediction of y:

$$H = \alpha_1 + \alpha_2 * x$$

$$y = \alpha_3 * H$$

→ Ideal solution:

$$\alpha_1 = b_0 = 2$$

$$\alpha_2 = b_1 = 4$$

$$\alpha_3 = 1$$

Let's compare  
with solution of  
NN ...

Fitted with **nnet()**  
function:


$$\alpha_1 = -0.19$$

$$\alpha_2 = 0.14$$


$$\alpha_3 = 115.93$$

!?

Look at  
predictions



Truth:  
 $b_0 = 2$   
 $b_1 = 4$






# Prediction of NN in more detail

- Univariate regression: Only one predictor (x)

- $y = b_0 + b_1x + \varepsilon \quad (\varepsilon \sim N(0, \sigma^2))$

- Applied Neural Network:



- Prediction of y:

$$H = \alpha_1 + \alpha_2 * x$$

$$y = \alpha_3 * H$$

→ Ideal solution:

$$\alpha_1 = b_0 = 2$$

$$\alpha_2 = b_1 = 4$$

$$\alpha_3 = 1$$

Let's compare  
with solution of  
NN ...

Fitted with **nnet()**  
function:


$$\alpha_1 = -0.19$$

$$\alpha_2 = 0.14$$


$$\alpha_3 = 115.93$$

!?


Look at  
predictions



Truth:  
 $b_0 = 2$   
 $b_1 = 4$



Zoom out ...




!?

# Prediction of NN in more detail

- Univariate regression: Only one predictor (x)

- $y = b_0 + b_1x + \varepsilon \quad (\varepsilon \sim N(0, \sigma^2))$

- Applied Neural Network:



- Prediction of y:

$$H = \alpha_1 + \alpha_2 * x$$

$$y = \alpha_3 * H$$

→ Ideal solution:

$$\alpha_1 = b_0 = 2$$

$$\alpha_2 = b_1 = 4$$

$$\alpha_3 = 1$$

Let's compare  
with solution of  
NN ...

Fitted with **nnet()**  
function:



$$\alpha_1 = -0.19$$

$$\alpha_2 = 0.14$$


$$\alpha_3 = 115.93$$

!?

Look at  
predictions




Zoom out ...



# Activation functions


- ▶ Activation functions transform the input of a neuron
- ▶ Generally, activation functions in the hidden layer are non-linear (e.g. logistic or step function)

▶ **Identity** function:




$$H_{output} = g(H_{input}) = H_{input}$$

▶ **Logistic** function:



$$H_{output} = g(H_{input}) = \frac{1}{1 + e^{-H_{input}}}$$

▶ **Step** function:






$$H_{output} = g(H_{input}) = \begin{cases} 1 & \text{if } H_{input} > 0 \\ 0 & \text{if } H_{input} \leq 0 \end{cases}$$

Activation function **g()**


$$H_{input} = \alpha_1 * X1 + \alpha_2 * X2$$

$$H_{output} = g(H_{input})$$




# Why non-linear activation functions?

- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )



Data




Neural Network


Prediction

# Why non-linear activation functions?


- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )




Data



Neural Network




Prediction






# Why non-linear activation functions?


- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )




Data



Neural Network




Prediction




Still just a line!

# Why non-linear activation functions?


- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )




Data



Neural Network




Prediction




# Why non-linear activation functions?


- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )




Data



Neural Network




Prediction




# Why non-linear activation functions?


- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )




Data



Neural Network




Prediction




Dashed lines show the individual contributions to the prediction of the three hidden nodes.

# Why non-linear activation functions?

- ▶ Only with non-linear activation functions can we model **non-linear patterns!**
- ▶ If only the identity function is used the output will **always be a linear function** (no matter how complex/deep the NN is)
- ▶ **Example:** Non-linear univariate regression
  - ▶  $y = x^2 + \varepsilon$  ( $\varepsilon \sim N(0, \sigma^2)$ )




Data




Show in 3d

Neural Network



Prediction



Dashed lines show the individual contributions to the prediction of the three hidden nodes.




# Neural Networks for classification

- ▶ The activation function and/or size in the output layer can be adapted to use NN for classification


## ▶ Regression

- ▶ Can use the identity function in the output layer




## ▶ Binary classification

- ▶ Can e.g. use the logistic function in the output layer
- ▶ The value of the output node represents the probability of  $y$  being equal to class 1



## ▶ Multiclass classification

- ▶ Could use  $k$  output nodes for  $k$  possible classes and e.g. a logistic function
- ▶ The  $k$  output nodes represent the probabilities that  $y$  is equal to a certain class
- ▶ Classically, the values of the  $k$  output nodes are forced to sum up to 1 (**softmax function**)




# Neural Networks for classification

- ▶ The activation function and/or size in the output layer can be adapted to use NN for classification


## ▶ Regression

- ▶ Can use the identity function in the output layer




## ▶ Binary classification

- ▶ Can e.g. use the logistic function in the output layer
- ▶ The value of the output node represents the probability of  $y$  being equal to class 1



## ▶ Multiclass classification



- ▶ Could use  $k$  output nodes for  $k$  possible classes and e.g. a logistic function
- ▶ The  $k$  output nodes represent the probabilities that  $y$  is equal to a certain class
- ▶ Classically, the values of the  $k$  output nodes are forced to sum up to 1 (**softmax function**)



# How are the weights found?

## ► "Training" of the Network

- Start with **random weights** (initialization)
- The predictions will be random as well (bad)
- Compare predictions with true y-values (classically using batches of training data)
  - Calculate how "wrong" the predictions were (**Loss-function**)
- Find out how the weights have to be shifted to improve the predictions (**backpropagation** algorithm to calculate gradients)
- Continue to update the weights until the algorithm converges
  - **Stochastic gradient descent**



# Single hidden layer NN in R

- ▶ Single hidden layer NNs can be fitted with **nnet** R-package

```

> library("nnet")
> nn_smk <- nnet(intention_to_smoke ~ friends_smoke + scale(alcohol_per_month) + scale(age),
  data = dat_smoking,
  size=2,
  decay=0,
  linout=FALSE,
  maxit=10000)

```

```
> summary(nn_smk)
```

a **3-2-1** network with 11 weights

options were - entropy fitting

```
b->h1 i1->h1 i2->h1 i3->h1
```

```
-3.88  9.32  2.74 -0.76
```

```
b->h2 i1->h2 i2->h2 i3->h2
```

```
9.16 -34.48 -7.01  2.28
```

```
b->o  h1->o  h2->o
```


```
-14.02  15.27  11.78
```

# Make predictions (here for training data):

```

> predict(nn_smk, newdata=dat_smoking,
  type='class')

```



# Single hidden layer NN in R

► Single hidden layer NNs can be fitted with **nnet** R-package

For neural networks it is usually advisable to standardize the predictors!

```
> library("nnet")
> nn_smk <- nnet(intention_to_smoke ~ friends_smoke + scale(alcohol_per_month) + scale(age),
  data = dat_smoking,
  size=2,
  decay=0,
  linout=FALSE,
  maxit=10000)
```

Number of nodes in hidden layer

"Weight decay" regularization (later)

Use identity activation function in output layer? (default is logistic)

Number of maximum iterations (how many iterations to adjust weights)

```
> summary(nn_smk)
```

a 3-2-1 network with 11 weights

options were - entropy fitting

b->h1 i1->h1 i2->h1 i3->h1

-3.88 9.32 2.74 -0.76

b->h2 i1->h2 i2->h2 i3->h2

9.16 -34.48 -7.01 2.28

b->o h1->o h2->o


-14.02 15.27 11.78

To predict labels use:

predict(..., type='class')

# Make predictions (here for training data):





```
> predict(nn_smk, newdata=dat_smoking,
  type='class')
```









# Parameter tuning in NNs

- ▶ With the **nnet()** function two main tuning parameter exists
  - ▶ "size" (number of neurons in hidden layer)
  - ▶ "decay" (Regularization factor using weight decay)
- ▶ The more neurons in the hidden layer the more complex patterns can be modeled
  - ▶ Danger of **under/over-fitting!**
- ▶ Overfitting on the example of the smoking data set:
  - ▶ Training error (miscl.rate) for different NN sizes



# Parameter tuning in NNs

- ▶ With the **nnet()** function two main tuning parameter exists
  - ▶ **"size"** (number of neurons in hidden layer)
  - ▶ **"decay"** (Regularization factor using weight decay)
- ▶ The more neurons in the hidden layer the more complex patterns can be modeled
  - ▶ Danger of **under/over-fitting!**
- ▶ Overfitting on the example of the smoking data set:
  - ▶ Training error (miscl.rate) for different NN sizes



```
> nnet(intention_to_smoke ~
friends_smoke +
scale(alc_hol_per_month) +
scale(age), data =
dat_smoking, size=1)
```

Training error:

22%

Under-fitting?

```
> nnet(intention_to_smoke ~
friends_smoke +
scale(alc_hol_per_month) +
scale(age), data =
dat_smoking, size=6)
```

Training error:

19%

Good?

```
> nnet(intention_to_smoke ~
friends_smoke +
scale(alc_hol_per_month) +
scale(age), data =
dat_smoking, size=20)
```

Training error:

16.5%

Over-fitting?

Can try to find  
optimum with  
**crossvalidation**

# Parameter tuning in NNs (caret)

- ▶ **Caret** is an R-package to automatically tune various machine learning models
- ▶ Allows the tuning of neural networks from **nnet** and **neuralnet** R-packages
- ▶ Have to define a **search grid** of parameter values
  - ▶ For each setting caret fits a NN and estimates the **test-error** using a resampling-based estimation (similar to crossvalidation)
  - ▶ The model with the lowest test-error is selected as the winner

```
> library("caret")  
### Create tuning grid:  
> t_grid <- expand.grid(size=c(1,5,10),  
                        decay=c(0, 0.5))  
  
> t_grid  
  size decay  
1    1  0.0  
2    5  0.0  
3   10  0.0  
4    1  0.5  
5    5  0.5  
6   10  0.5  
  
### Tune the model ("train" function):  
> set.seed(288)  
> tune_caret <- train(intention_to_smoke ~ .,  
                      data=dat_smoking, method='nnet', tuneGrid=t_grid,  
                      maxit=10000, linout=FALSE, preProcess=c("center","scale"))
```

# Parameter tuning in NNs (caret)

- ▶ **Caret** is an R-package to automatically tune various machine learning models
- ▶ Allows the tuning of neural networks from **nnet** and **neuralnet** R-packages
- ▶ Have to define a **search grid** of parameter values
  - ▶ For each setting caret fits a NN and estimates the **test-error** using a resampling-based estimation (similar to crossvalidation)
  - ▶ The model with the lowest test-error is selected as the winner

```
> tune_caret <- train(intention_to_smoke ~ .,  
data=dat_smoking, method='nnet', tuneGrid=t_grid,  
maxit=10000, linout=FALSE, preProcess=c("center","scale"))
```

```
> tune_caret  
Neural Network  
200 samples  
4 predictor  
2 classes: 'no', 'yes'
```

Pre-processing: centered (4), scaled (4)

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 200, 200, 200, 200, 200, 200, ...

Resampling results across tuning parameters:

size	decay	Accuracy	Kappa
1	0.0	0.7257813	0.4456960
1	0.5	0.7231326	0.4456194
5	0.0	0.6825685	0.3546761
5	0.5	0.7219893	0.4428660
10	0.0	0.6467391	0.2792476
10	0.5	0.7207855	0.4403016

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were size = 1 and decay = 0.

# Parameter tuning in NNs (caret)

- ▶ **Caret** is an R-package to automatically tune various machine learning models
- ▶ Allows the tuning of neural networks from **nnet** and **neuralnet** R-packages
- ▶ Have to define a **search grid** of parameter values
  - ▶ For each setting caret fits a NN and estimates the **test-error** using a resampling-based estimation (similar to crossvalidation)
  - ▶ The model with the lowest test-error is selected as the winner
  - ▶ The (hyper) parameters of the winner model are used to fit a final model to the complete data (this will be the model used for prediction)
  - ▶ Can use the predict function with the caret object: `predict(tune_caret)`

```
> tune_caret <- train(intention_to_smoke ~ .,
  data=dat_smoking, method='nnet', tuneGrid=t_grid,
  maxit=10000, linout=FALSE, preProcess=c("center","scale"))
```

```
> tune_caret
Neural Network
200 samples
  4 predictor
  2 classes: 'no', 'yes'
```

Pre-processing: centered (4), scaled (4)

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 200, 200, 200, 200, 200, 200, ...

Resampling results across tuning parameters:

size	decay	Accuracy	Kappa
1	0.0	0.7257813	0.4456960
1	0.5	0.7231326	0.4456194
5	0.0	0.6825685	0.3546761
5	0.5	0.7219893	0.4428660
10	0.0	0.6467391	0.2792476
10	0.5	0.7207855	0.4403016

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were size = 1 and decay = 0.

# Parameter tuning in NNs (caret)

- ▶ **Caret** is an R-package to automatically tune various machine learning models
- ▶ Allows the tuning of neural networks from **nnet** and **neuralnet** R-packages
- ▶ Have to define a **search grid** of parameter values
  - ▶ For each setting caret fits a NN and estimates the **test-error** using a resampling-based estimation (similar to crossvalidation)
  - ▶ The model with the lowest test-error is selected as the winner
  - ▶ The (hyper) parameters of the winner model are used to fit a final model to the complete data (this will be the model used for prediction)
  - ▶ Can use the predict function with the caret object: `predict(tune_caret)`

```
> tune_caret <- train(intention_to_smoke ~ .,
  data=dat_smoking, method='nnet', tuneGrid=t_grid,
  maxit=10000, linout=FALSE, preProcess=c("center","scale"))
```

```
> tune_caret
Neural Network
200 samples
  4 predictor
  2 classes: 'no', 'yes'
```

Pre-processing: centered (4), scaled (4)

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 200, 200, 200, 200, 200, 200, ...

Resampling results across tuning parameters:

size	decay	Accuracy	Kappa
1	0.0	0.7257813	0.4456960
1	0.5	0.7231326	0.4456194
5	0.0	0.6825685	0.3546761
5	0.5	0.7219893	0.4428660
10	0.0	0.6467391	0.2792476
10	0.5	0.7207855	0.4403016

Because we evaluated multiple models, the performance estimate of the winner configuration can be too optimistic. Therefore, people often use an additional separate data set to get a final performance estimate after model selection.

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were size = 1 and decay = 0.

# Picture recognition with NNs

- ▶ So far in supervised learning: Data is always a **table** with **predictors** and **target variable**
- ▶ What if we want to predict the content shown on a **picture**?
  - ▶ A picture is no different! We can translate it into a row of values where each value represents the shade of one **pixel**.


- ▶ Example: Predict/recognize **handwritten digits**

- ▶ Exemplary data table of pictures with 30 x 30 pixels:


	Pix1	Pix2	Pix3	...	Pix900	Digit
Picture 1	20	24	60	...	44	<b>0</b>
Picture 2	10	94	160	...	244	<b>7</b>
Picture 3	220	89	143	...	134	<b>3</b>
Picture 4	12	123	70	...	230	<b>5</b>
...	...	...	...	...	...	...

Grayscale values of all 900 pixels on a picture

**Target variable**



The shading of each pixel corresponds to a "**grayscale**" value ranging from 0 (white) to 255 (black).






# Picture recognition with NNs

- ▶ Grayscale images are normally stored as tables with rows and columns indicating the **pixel positions**

- ▶ Table storing a grayscale picture with 4 x 4 pixels:


100	0	0	255
0	0	255	0
0	255	0	0
255	0	0	100

Is the image



- ▶ A collection of multiple images is normally stored as a **3-dimensional array**, which is like a cube with pictures "stacked" as layers

- ▶ Exemplary array storing four 4x4 pixel images:



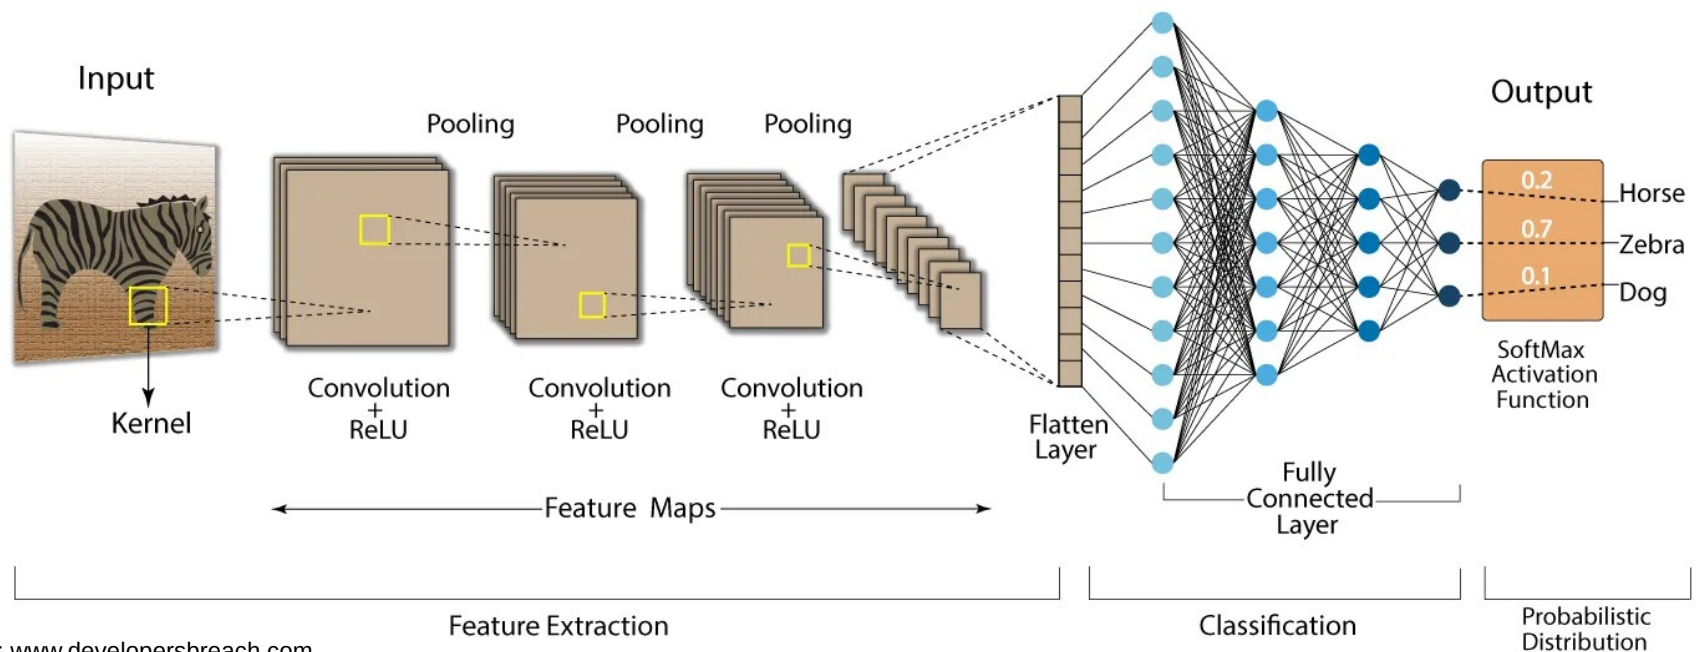
We have to **"flatten"** the images so that we can feed them to a classifier like a NN.

	Pix1	Pix2	Pix3	...	Pix16
Picture 1	20	24	60	...	44
Picture 2	10	94	160	...	244
Picture 3	220	89	143	...	134
Picture 4	12	123	70	...	230

(table-format we need)

# Picture recognition with NNs

- ▶ Feeding a "flattened" image into a classification NN works for smaller picture sizes
- ▶ For large pictures, however, the NN becomes too complex and complicated (millions of parameters)
- ▶ **Convolutional NNs** try to more efficiently capture the spatial dependencies in a picture by reducing the image to a set of (hopefully) relevant features (**feature extraction**)
  - ▶ The extracted features are subsequently fed into a classification NN



# Deep learning with R (book)

- ▶ Good introduction to deep neural networks with R (by François Chollet and J.J. Allaire)

